

CVaR-Based Portfolio Optimisation for a Malta-Accessible UCITS ETF Basket

Abdirizak Hassan Osman
Institute of Information & Communication Technology
Malta College of Arts, Science & Technology
Corradino Hill
Paola PLA 9032
abdirizak.hassan.e21736@mcast.edu.mt

Abstract—This paper offers some additional guidelines for MCAST IICT 2nd year B.Sc. students on research paper writing. The abstract is the first paragraph that any researcher reads and thus you need to capture the essence of your research. Dedicate 1 sentence for the theme (subject matter), another for your project aim, one for the proposed solution, then for the general outcome of this project (positive/negative result/outcome), another for the technology used (Windows/Linux, Android/iOS, C#/Java, Unity/Unreal), finally for any technique used (Neural Networks, Pearl Noise, Market Basket Analysis, K-Means, HMM, Augmented Reality). Following is an example: *In this research we are tackling the automatic annotation of cast members and use of key props within movies. Movies tend to gather a huge fan gathering who dedicate a lot of time and resources for documenting the content, thus tools to automate processes such as the timeline when cast members appear, presents of props such as weapons and other forms of annotations are desired. In this research we propose the use of image processing techniques, namely Principal Component Analysis (PCA) and Convolutional Neural Networks (CNN) for the automatic annotation of actors' timeline and presence of weapons in the popular Tv series Game of Thrones (GoT). Our proposed solution managed to detect the targets in 85% of the cases and we identify situations where this is challenging and recommend future research directions.*

Index Terms—MCAST, IICT, L^AT_EX, Project, Paper

I. INTRODUCTION

Recent market crises, including the COVID-19 shock, have renewed criticism of classical portfolio theory. Although the mean-variance (MV) framework remains foundational, variance treats upside and downside deviations symmetrically and therefore does not isolate severe losses in the way downside-focused investors typically require [1], [2]. For this reason, portfolio research has increasingly adopted Conditional Value-at-Risk (CVaR), which estimates average losses beyond the VaR threshold and can be incorporated into tractable optimisation formulations [2], [3]. Exchange-traded funds (ETFs) provide a practical setting for such comparison because prior studies have evaluated CVaR-based, MV, and equal-weight approaches using diversified ETF universes and out-of-sample backtesting [3].

Despite this relevance, much of the recent CVaR literature is embedded in advanced settings such as evolutionary algorithms, copula-based dependence modelling, bilevel multi-market optimisation, deep reinforcement learning, and distributionally robust optimisation [2]–[6]. These studies are

valuable, but they make it harder to judge whether a simpler and more transparent CVaR-minimisation framework, implemented under standard long-only constraints, already provides meaningful downside-risk benefits. In addition, while Malta has an established collective investment scheme environment and the Maltese Assets Fund illustrates the local relevance of diversified fund products [7], [8], there is limited evidence on CVaR-based optimisation in a Malta-accessible UCITS ETF context. The present study is therefore positioned at the intersection of quantitative portfolio theory and applied investment management, addressing a contextual and empirical gap by testing whether a Malta-accessible UCITS ETF basket can serve as a credible benchmark setting for comparing CVaR, Equal-Weight, and Minimum-Variance strategies.

The purpose of this quantitative backtesting study is to evaluate whether CVaR-minimisation provides stronger downside-risk control than Equal-Weight (1/N) and Minimum-Variance portfolio construction in a Malta-accessible UCITS ETF basket over 2016–2025. The independent variable is the optimisation strategy, operationalised through three portfolio rules: Equal-Weight, Minimum-Variance, and CVaR-minimisation. The dependent variables are the observed portfolio outcomes, namely CAGR, volatility, Sharpe ratio, Sortino ratio, VaR, CVaR, and maximum drawdown. To preserve comparability, the ETF universe, estimation window length, rebalancing frequency, and long-only fully invested constraints are treated as control variables. The study aims to implement the three strategies under identical conditions using rolling out-of-sample backtesting and to compare their performance on both return-based and downside-risk criteria.

Accordingly, the main hypothesis states that a CVaR-minimised portfolio will achieve lower CVaR and lower maximum drawdown than Equal-Weight and Minimum-Variance portfolios without material deterioration in CAGR or Sharpe ratio. The null hypothesis states that no meaningful difference will exist among the three strategies in CVaR, maximum drawdown, or risk-adjusted performance. As shown in Fig. 1, the present study adopts a positivist philosophy and a deductive approach. It employs a controlled experimental backtesting strategy using secondary historical ETF data, a mono-method quantitative design, and a longitudinal time horizon. The techniques and procedures consist of secondary ETF data

collection, rolling out-of-sample backtesting, and quantitative performance and downside-risk analysis. Specifically, portfolio outcomes are treated as objectively measurable through predefined financial metrics. A deductive approach is adopted because the study begins from established theory on the limitations of variance as a risk proxy and tests whether a downside-risk measure produces superior empirical outcomes [1], [2]. Since performance is evaluated across repeated windows over 2016–2025 rather than at a single point in time, the time horizon is longitudinal.




Fig. 1. Research onion for the present study, adapted from Fig. 4.1 in Saunders, Lewis, and Thornhill [9].

II. LITERATURE REVIEW

Section II reviews the research methodologies used in six peer-reviewed journal articles [1]– [6]. Non-academic and institutional sources are used only for contextual background and are excluded from the methodological comparison.

Fig. 2 organises the reviewed methodological literature into data employed, process and candidate techniques, and evaluation and testing strategies. The map shows that prior studies mainly rely on secondary financial data and backtesting-based evaluation, but differ substantially in portfolio-construction complexity. The final branch highlights the gap addressed by the present study, namely a simpler and more transparent comparison of Equal-Weight, Minimum-Variance, and CVaR-minimisation strategies in a Malta-accessible UCITS ETF setting under identical long-only conditions.

A. Overview of reviewed methodologies

A consistent methodological pattern across the reviewed studies is the use of a positivist, quantitative design based on secondary financial data. Rather than using interviews or surveys, the reviewed papers typically rely on historical price or return series and assess portfolio-construction rules through computational testing. This broad methodological direction aligns closely with the present study, which also examines objectively measurable portfolio outcomes under controlled

backtesting conditions. Accordingly, the literature map for the present study is organised around data, process, and evaluation, as illustrated in Fig. 2.

B. Data employed in prior studies

The reviewed literature shows that portfolio-optimisation research commonly uses historical market data covering several years so that both normal and stressed market conditions are represented. Muñoz and Cifuentes [1], for example, analyse historical returns from 2004 to 2023 and deliberately include major crisis periods to test the stability of mean-variance and CVaR formulations. Their use of rolling windows is methodologically valuable because it allows repeated comparison across time rather than relying on a single in-sample period.

Teplova et al. [3] also use a long historical dataset, drawing on daily returns from international ETFs between 2012 and 2022. This is particularly relevant because ETFs are also the asset class used in the present study. Their data design supports realistic comparison between portfolio rules while also enabling out-of-sample testing.

By contrast, Nasini et al. [4] use a much larger firm-level dataset across thousands of U.S. listed companies, while Cui et al. [5] focus on cryptocurrency returns. Kobayashi et al. [6] also work with large asset universes in a robust optimisation setting. These studies show that broader datasets are possible, but they also increase modelling and computational complexity. For the current research, a Malta-accessible UCITS ETF basket is more appropriate because it remains sufficiently diversified while preserving transparency, feasibility, and direct relevance to the stated research context.

Overall, the literature suggests that the most suitable data strategy for the present study is a secondary dataset of historical ETF prices and returns, covering a sufficiently long period to include both calm and turbulent market conditions, while keeping the investable universe constant across all tested strategies.

C. Portfolio-construction process and candidate techniques

The reviewed studies show that the main methodological distinction lies in the portfolio-construction process. Some studies compare classical benchmark approaches, while others embed CVaR within more advanced optimisation frameworks.

Muñoz and Cifuentes [1] provide the clearest benchmark-comparison methodology. They compare mean-variance and CVaR formulations under the same target-return structure, which isolates the effect of the risk metric itself. This is highly relevant to the present study because it demonstrates how a fair experimental comparison can be achieved when constraints and data inputs are held constant.

Teplova et al. [3] extend the process stage through copula modelling and a Black-Litterman mean-CVaR framework. Their methodology is rigorous, but it depends on parameter calibration, investor-view modelling, and dependence-structure estimation. Kaucic et al. [2] instead treat the process stage as a multi-objective evolutionary search problem, comparing variants of NSGA-II and SPEA 2 under downside-risk measures




Fig. 2. Methodological literature map for CVaR-based portfolio optimisation.

such as semi-variance and CVaR. Nasini et al. [4] introduce a bilevel multi-market formulation, while Kobayashi et al. [6] use distributionally robust optimisation with cardinality constraints. Cui et al. [5] move further still by integrating deep reinforcement learning into a mean-CVaR framework.

These advanced methodologies are useful as evidence that CVaR is a mature and flexible risk measure, but they are not ideal templates for the present project. The current study is not intended to propose a highly specialised optimisation engine. Instead, it aims to compare three portfolio-construction rules under identical conditions: Equal-Weight, Minimum-Variance, and CVaR-minimisation. This choice is methodologically stronger for the present study because it prioritises interpretability, control, and reproducibility over algorithmic novelty.

D. Evaluation and testing strategies

The literature also shows a clear methodological pattern in the evaluation stage. Studies rarely stop at in-sample optimisation. Instead, they use out-of-sample testing and compare strategies through financial performance and risk metrics.

Muñoz and Cifuentes [1] use rolling windows and assess the similarity or difference between resulting portfolios under different risk formulations. Teplova et al. [3] evaluate portfolio performance against benchmark strategies and report risk-adjusted outcomes as well as tail-risk behaviour. Kobayashi et al. [6] also use out-of-sample investment performance measures, while Kaucic et al. [2] assess optimisation quality through Pareto-front indicators in addition to portfolio risk measures.

Across the reviewed studies, the most relevant evaluation measures for the present research are those that capture both return and downside risk. In particular, CAGR, volatility, Sharpe ratio, Sortino ratio, VaR, CVaR, and maximum draw-down appear suitable because they allow direct comparison

between standard performance and tail-risk protection. This is important because the study hypothesis does not claim that CVaR-minimisation will maximise returns; rather, it tests whether CVaR-minimisation improves downside-risk control without materially worsening return-based outcomes.

E. Comparative synthesis and methodological gap

Taken together, the reviewed literature supports three main conclusions. First, the dominant research tradition in this area is quantitative, secondary-data, and backtesting-based. This strongly supports the philosophical and methodological position already established in Section I and illustrated in Fig. 1.

Second, the strongest comparative designs are those that keep the dataset, constraints, and testing framework constant while varying only the portfolio rule. This is the methodological logic most clearly visible in Muñoz and Cifuentes [1], and it is the one most appropriate for the present study.

Third, although advanced methods such as copulas, bilevel programming, evolutionary optimisation, robust optimisation, and deep reinforcement learning increase technical novelty, they also introduce greater complexity and reduce transparency. For this research a simpler controlled comparison of benchmark strategies is methodologically more appropriate.

Finally, across the reviewed studies, the main independent element is the portfolio-construction or optimisation method, while the dependent outcomes are portfolio performance and downside-risk measures assessed through out-of-sample testing. Common testing strategies include rolling-window backtesting, benchmark comparison, and risk-adjusted performance evaluation. However, despite this methodological maturity, limited evidence compares CVaR-minimisation, minimum-variance, and equal-weight strategies in a simple, transparent, Malta-accessible UCITS ETF setting under identical long-only backtesting conditions.

III. RESEARCH METHODOLOGY

Portfolio construction depends heavily on how risk is defined and measured. Although mean-variance optimisation remains a classical benchmark, variance treats upside and downside deviations symmetrically and may therefore understate the practical importance of severe losses during turbulent market conditions [1], [3]. By contrast, Conditional Value-at-Risk (CVaR) focuses on the average loss beyond the VaR threshold and has been widely adopted in portfolio research as a downside-risk measure that is more sensitive to tail losses [1]–[3]. However, much of the recent literature applies CVaR within complex optimisation settings such as copula-based models, evolutionary algorithms, robust optimisation, or deep reinforcement learning [2]–[6]. This makes it less clear whether a simpler and more transparent CVaR-minimisation framework can already provide meaningful downside-risk benefits when compared with standard benchmark strategies.

In view of this problem, the present study examined whether portfolio construction based on CVaR-minimisation offers stronger downside-risk protection than Equal-Weight and Minimum-Variance strategies in a Malta-accessible UCITS ETF basket. The comparison is conducted under identical long-only and fully invested conditions using secondary historical ETF data and rolling out-of-sample backtesting. This design allows the effect of the portfolio-construction rule to be examined more directly, while keeping the dataset, constraints, and testing framework consistent across all strategies [5], [6].

The main hypothesis of this study states that, when applied to the same Malta-accessible UCITS ETF basket under identical rolling out-of-sample backtesting conditions, a CVaR-minimised portfolio will achieve lower CVaR and lower maximum drawdown than Equal-Weight and Minimum-Variance portfolios, without material deterioration in CAGR or Sharpe ratio. The corresponding null hypothesis states that no meaningful difference will exist among the three strategies in terms of CVaR, maximum drawdown, or risk-adjusted performance. This hypothesis is testable through empirical investigation because the independent variable is the portfolio-construction rule, while the dependent variables are the observed portfolio outcomes measured through predefined financial metrics [1], [5].

In light of the above, the aim of this research is to assess whether a simpler and more transparent CVaR-minimisation framework can provide stronger downside-risk protection than Equal-Weight and Minimum-Variance portfolio construction in a Malta-accessible UCITS ETF basket. More specifically, the study seeks to implement the three strategies under identical long-only and fully invested conditions and to evaluate their out-of-sample performance over the 2016–2025 period. The expected outcome is a clear empirical comparison showing whether CVaR-minimisation improves downside-risk measures such as CVaR and maximum drawdown, while preserving broadly competitive return and risk-adjusted performance.

The research objectives were as follows:

- 1) Construct a consistent historical dataset for the selected

Malta-accessible UCITS ETF basket and prepare it for portfolio backtesting.

- 2) Implement Equal-Weight, Minimum-Variance, and CVaR-minimisation portfolio-construction rules under identical investment constraints.
- 3) Conduct rolling out-of-sample backtesting to evaluate how the three strategies perform across different market conditions over 2016–2025.
- 4) Compare the strategies using return, volatility, risk-adjusted, and downside-risk measures, namely CAGR, volatility, Sharpe ratio, Sortino ratio, VaR, CVaR, and maximum drawdown.
- 5) Assess whether CVaR-minimisation provides materially stronger downside-risk control than the benchmark strategies without significant deterioration in return-based performance.

The research questions addressed in this study are organised around the data, portfolio-construction process, and evaluation stages identified in Section II. Since the present research does not propose a new complex optimisation engine, the questions are designed to examine whether a transparent benchmark comparison based on secondary ETF data can produce meaningful evidence on downside-risk control.

The research questions are as follows:

- 1) Historical dataset: What historical ETF dataset and investable UCITS ETF basket are most suitable for a controlled comparison of Equal-Weight, Minimum-Variance, and CVaR-minimisation strategies in a Malta-accessible setting?
- 2) process / candidate techniques: Under identical long-only and fully invested constraints, how does CVaR-minimisation compare with Equal-Weight and Minimum-Variance as a portfolio-construction approach?
- 3) evaluation and testing: How does the proposed CVaR-minimisation strategy compare with the benchmark strategies in rolling out-of-sample backtesting when evaluated using CAGR, volatility, Sharpe ratio, Sortino ratio, VaR, CVaR, and maximum drawdown?

A. Research Pipeline

To answer the above research questions, the study was structured around the research pipeline illustrated in Fig. 3. The first stage consisted of defining the six-ETF UCITS basket and collecting the historical adjusted closing prices required for the backtest. The second stage involved aligning prices across common trading dates, checking completeness, handling missing observations, and deriving periodic returns from the cleaned series. The third stage consisted of constructing Equal-Weight, Minimum-Variance, and CVaR-minimisation portfolios under identical long-only and fully invested constraints. The fourth stage involved rolling out-of-sample backtesting, where portfolio weights were estimated from the preceding in-sample window and then applied to the subsequent out-of-sample period until the end of 2025. The final stage consisted of evaluating realised portfolio outcomes

using CAGR, annualised volatility, Sharpe ratio, Sortino ratio, VaR, CVaR, and maximum drawdown, and comparing the CVaR strategy against the benchmark portfolios.




Fig. 3. Fig. 3. Research pipeline for CVaR-based portfolio optimisation in a Malta-accessible UCITS ETF basket.

B. Research Approach and Experimental Variables

Variable type	Variable	Operationalisation
Independent variable	Portfolio-construction strategy	Equal-Weight, Minimum-Variance, CVaR-minimisation.
Dependent variables	Portfolio outcomes	CAGR, annualised volatility, Sharpe ratio, Sortino ratio, VaR, CVaR, maximum draw-down.
Control variables	ETF universe	Fixed six-ETF Malta-accessible UCITS basket.
Control variables	Data treatment	Adjusted closing prices, common trading dates, identical return-construction procedure.
Control variables	Investment constraints	Long-only, fully invested.
Control variables	Backtesting settings	Same rolling estimation logic, rebalancing structure, and evaluation period for all strategies.

The present study adopted a positivist philosophy, a deductive approach, and a mono-method quantitative design. This choice was appropriate because the research aimed to test whether changing the portfolio-construction rule produced measurable differences in objectively observed portfolio outcomes. More specifically, the study employed a controlled experimental backtesting strategy using secondary historical ETF data. This design was methodologically suitable because the ETF universe, data treatment, and investment constraints were held constant, while the portfolio-construction rule was varied across Equal-Weight, Minimum-Variance, and CVaR-minimisation. In this way, the study was designed to isolate the effect of the optimisation approach as clearly as possible.

The independent variable in the study was the portfolio-construction strategy, operationalised through three portfolio rules: Equal-Weight, Minimum-Variance, and CVaR-minimisation. The dependent variables were the observed portfolio outcomes, namely CAGR, annualised volatility, Sharpe ratio, Sortino ratio, VaR, CVaR, and maximum drawdown. To preserve comparability, the ETF universe, adjusted-price treatment, common-date alignment, long-only fully invested constraints, and the same rolling out-of-sample testing framework were treated as control variables. Since the study evaluated repeated portfolio outcomes over time rather than collecting subjective human responses, the design was consistent with a quantitative backtesting methodology.

C. Dataset and Data Preparation

The dataset used in this study consisted of secondary historical ETF market-price data for a fixed basket of six Malta-accessible UCITS ETFs, representing four equity sleeves and two government-bond sleeves. The selected basket comprised iShares Core S&P 500 UCITS ETF USD (Acc), iShares Core MSCI Europe UCITS ETF EUR (Acc), iShares Core MSCI Japan IMI UCITS ETF, iShares Core MSCI Emerging Markets IMI UCITS ETF (Acc), iShares USD Treasury Bond

7–10yr UCITS ETF (Acc), and iShares Euro Government Bond 7–10yr UCITS ETF (Acc). ISINs were treated as the master identifiers for the basket, while EUR was used as the reporting currency for analysis.

Historical adjusted closing prices were downloaded as an exchange market-price dataset through the yfinance interface and stored locally for reproducibility. Adjusted prices were used because they provided a more consistent basis for return calculation than unadjusted closing prices. The raw data were then aligned across common trading dates, checked for completeness, and cleaned to handle missing observations before periodic returns were computed from the aligned series. This preparation stage was necessary to ensure that all three portfolio-construction rules were tested on the same underlying data under identical conditions.

D. Portfolio-Construction and Rolling Out-of-Sample Backtesting

Three portfolio-construction strategies were implemented in the study. Equal-Weight served as a transparent benchmark in which capital was allocated equally across all ETFs. Minimum-Variance represented the classical variance-based optimisation approach, while CVaR-minimisation represented the downside-risk approach under investigation. All three strategies were implemented under shared investment constraints so that no short selling was permitted and the full portfolio weight remained invested at each rebalance point. This ensured that the comparison remained controlled, interpretable, and directly aligned with the hypothesis.

To evaluate these strategies, the study employed rolling out-of-sample backtesting over the 2016–2025 period. At each rebalance date, portfolio weights were estimated from the preceding in-sample window and then applied to the subsequent out-of-sample holding period. The window was then rolled forward and the procedure repeated until the end of the sample. This design was preferred to a single in-sample optimisation because it reduced look-ahead bias and enabled repeated testing across different market conditions. In methodological terms, this rolling structure was also consistent with the evaluation patterns identified in the reviewed portfolio-optimisation literature.

A separate pilot study was not required because the research did not involve human participants, questionnaires, or competing candidate systems that first needed to be screened before main testing. Instead, preliminary implementation checks were carried out to verify price alignment, return calculation, optimisation constraints, and output consistency before running the full backtest. This was sufficient for establishing a stable quantitative workflow without introducing an unnecessary additional experimental phase.

E. Performance Evaluation

The realised out-of-sample portfolio results were evaluated using both return-oriented and downside-risk metrics. Return and risk-adjusted performance were assessed using CAGR, annualised volatility, Sharpe ratio, and Sortino ratio.

Downside-risk behaviour was assessed using VaR, CVaR, and maximum drawdown. These measures were selected because the hypothesis did not claim that CVaR-minimisation would maximise return; rather, it tested whether downside-risk control could be improved without materially worsening return-based performance.

The analysis was primarily comparative. The EW, MV, and CVaR portfolios were compared under the same dataset, constraints, and backtesting logic, and the resulting metrics were interpreted relative to the hypothesis and to the benchmark-comparison designs discussed in Section II. This approach was consistent with the assignment guidance that findings should later be discussed honestly and compared with existing research rather than being presented as a perfect or universal solution.

F. Validity, Reliability, Generalisability, Limitations, and Ethical Considerations

The validity of the study was strengthened by the controlled design. All three portfolio rules were tested on the same ETF basket, over the same time horizon, using the same adjusted-price treatment, date alignment, and investment constraints. The main intentional source of variation was therefore the portfolio-construction rule itself. Reliability was supported through the use of fixed procedures for data preparation, return calculation, portfolio construction, and backtesting, together with local storage of the downloaded data for reproducibility. However, the generalisability of the findings remained bounded by the specific six-ETF UCITS basket, the 2016–2025 period, and the long-only fully invested framework used in the study.

The study also had several limitations. First, it relied on a relatively small and intentionally transparent ETF universe rather than a large institutional asset set. Second, portfolio outcomes depended on historical market behaviour and therefore could not guarantee future performance. Third, the backtest was based on market-price data and did not attempt to model every real-world friction in depth, such as taxes or all trading frictions. These limitations do not invalidate the study, but they do define the scope within which its conclusions should be interpreted.

Finally, the study raised minimal ethical risk because it used secondary historical financial data and did not involve human participants, interviews, surveys, or personal data. The main ethical obligations were therefore academic honesty, correct citation, transparent reporting of the method, and accurate acknowledgment of the data source and implementation choices. This is consistent with the wider MCAST guidance that methodology should be rigorous, clearly explained, and properly referenced.

IV. FINDINGS & DISCUSSION OF RESULTS

This section is a delicate section and you are encouraged to be as honest and open as possible. The aim is not to show that you got a perfect solution to a long existing problem. Trying to state that in a few months you developed a perfect

solution with 100% is not convincing and raises doubts. It is recommended that you document your findings in every step of your research pipeline, highlighting your observations and decisions taken. Present your results and very importantly compare with existing research which you documented in Section II.

The focus of the Project module is for you to delve into an area that exposes you to new technologies and offers you an opportunity to be critical of your work. So you are expected to document where your solution/research worked and where it did not. Reflect and document reasons why the solution/research did not perform as expected and propose ways of addressing this. From these observations you will produce new research questions in the next section. Consider the research in “Brand usage detection within audio streams”, where certain key terms were searched within videos, the results of which are documented in Table II. You will notice that the terms “Peppa” and “Sushi” were the least recognised terms even by the best transcribers. Upon investigation we determined that “Peppa” was not recognised cause of voice morphing to create childish voices in the cartoon video, whilst “Sushi” was pronounced by a Japanese person speaking English. So the research student decided to focus his dissertation research on how to create a system that is able to recognise heavy accents to automate the configuration of a transcriber, in this case to cater for English spoken by a Japanese person, which accent is very different from an Indian accent, British accent or Italian accent, just to name a few.

TABLE II
RECALL RESULTS

#	Term	Google Cloud	Google Speech	Sphinx CMU
1	Peppa	27%	33%	0%
2	Peppa	33%	22%	0%
3	Apple	96%	92%	79%
4	Galaxy	100%	100%	100%
5	Galaxy	95%	95%	80%
6	Sushi	75%	35%	0%
	Average	71%	62%	43%

V. CONCLUSION

So the conclusion is most probably the second section that a reader would use to consider reading in full your research. Thus it is important to highlight the essence of your research. The recommended approach is to answer your research methodology. Start by answering your research questions, then stating to what degree did this research achieve its aim and objectives, highlighting potential causes for not being able to do so at a desired level, such as time, or other circumstances. Consider the following: *A student was due to research the use of MCAST computers during out-of-office hours to offer a private cloud computing service for research, similar to Google Cloud but free for MCAST students. Due to the lockdown by COVID-19 pandemic we could not continue on the original planned research objective and had to adapt.*

The final and most important part of the conclusion are your recommendations for future research, not necessarily for

yourself (referring to what you plan to do in your dissertation or beyond), but also to other future researchers who might consider doing similar work to yours. The recommendations you provide here will set such prospective researchers on a better track/direction thanks to your experience.

APPENDIX A SUPPORTING MATERIAL

You can add screen shots and statistics. Stick to essential information.

ACKNOWLEDGEMENT

You can dedicate this section for special assistance that you were given by a 3rd party. Not your mentor, relatives or questionnaire participants.

REFERENCES

- [1] I. Muñoz and A. Cifuentes, “Portfolio optimization: Do risk metrics matter?” *Finance Research Open*, vol. 2, no. 1, p. 100098, Mar. 2026.
- [2] M. Kaucic, M. Moradi, and M. Mirzazadeh, “Portfolio optimization by improved nsga-ii and spea 2 based on different risk measures,” *Financial Innovation*, vol. 5, no. 1, p. 26, Dec. 2019.
- [3] T. Teplova, M. Evgeniia, Q. Munir, and N. Pivnitskaya, “Black-litterman model with copula-based views in mean-cvar portfolio optimization framework with weight constraints,” *Economic Change and Restructuring*, vol. 56, no. 1, pp. 515–535, Feb. 2023.
- [4] S. Nasini, M. Labbé, and L. Brotcorne, “Multi-market portfolio optimization with conditional value at risk,” *European Journal of Operational Research*, vol. 300, no. 1, pp. 350–365, Jul. 2022.
- [5] T. Cui, S. Ding, H. Jin, and Y. Zhang, “Portfolio constructions in cryptocurrency market: A cvar-based deep reinforcement learning approach,” *Economic Modelling*, vol. 119, p. 106078, Feb. 2023.
- [6] K. Kobayashi, Y. Takano, and K. Nakata, “Cardinality-constrained distributionally robust portfolio optimization,” *European Journal of Operational Research*, vol. 309, no. 3, pp. 1173–1182, Sep. 2023.
- [7] HSBC Global Asset Management (Malta) Ltd., “Maltese assets fund: Fund information sheet, income class,” Fund information sheet, Jan. 2026.
- [8] Malta Financial Services Authority, “Net asset value (eur) - locally-based collective investment schemes,” Malta Financial Services Authority, Tech. Rep., Mar. 2025, statistical report. [Online]. Available: <https://www.mfsa.mt/wp-content/uploads/2025/05/Net-Asset-Value-Locally-Based-Collective-Investment-Schemes-March-2025.pdf>
- [9] M. Saunders, P. Lewis, and A. Thornhill, *Research Methods for Business Students*, 4th ed. Harlow, U.K.: Financial Times Prentice Hall, 2007.